

[Home](#)[Project 2](#)

# Convolution and Filtering

CS180 • Fall 2025

[Back to Home](#) ›

## 1A – Convolution Filters

This project studies linear filtering and multi-scale image representations through a sequence of controlled experiments. In Part 1, I implement 2-D convolution from first principles, estimate spatial derivatives via finite differences, and replace naive differencing with derivative-of-Gaussian (DoG) filters to suppress noise. In Part 2, I treat filtering as a tool for manipulating frequency content: unsharp masking increases the energy of high spatial frequencies to “sharpen” imagery; hybrid images juxtapose low-frequency content from one image with high-frequency content from another; and Gaussian/Laplacian stacks enable multi-resolution blending that hides seams across scales. All results are generated by my own code; third-party routines are used only as references for verification and timing.

**Dataset:** Box Filter ▾ **k = 3**  ☒ From scratch

☐ SciPy





**Box Filter — From scratch (k=3)**

I implemented convolution in `proj2/conv.py` in two forms. The pedagogical baseline, `conv2d_four_loops`, explicitly flips the kernel and accumulates the product over a zero-padded neighborhood using four Python loops (two for spatial coordinates, two for kernel indices). The optimized variant, `conv2d_two_loops`, maintains identical semantics but reduces interpreter overhead by vectorizing the inner multiply-accumulate across complete row (or column) slices, leaving only two Python loops over output coordinates. For validation and runtime comparison I use `conv2d_scipy`, a thin wrapper around `scipy.signal.convolve2d` with *same* output size and *fill* (zero) boundary conditions. The box filter used in this section is created by `proj2/filters.box_filter(k)`, which returns a constant kernel whose entries sum to one, preserving DC gain. Visual and numerical comparisons confirm that both custom implementations match SciPy up to floating-point tolerance, while the two-loop version is substantially faster than four loops for moderate kernel sizes.

## 1B – Finite Difference Operator

To estimate spatial derivatives, I convolve a grayscale image with the forward-difference stencils  $\mathbf{dx} = [-1, 1]$  and  $\mathbf{dy} = \mathbf{dx}^T$  using `conv2d_scipy`. The partials  $I_x$  and  $I_y$  are then combined into a gradient magnitude image  $|\nabla I| = \sqrt{I_x^2 + I_y^2}$  (implemented in `proj2/edges.py`). Binarized “edge maps” are obtained by thresholding  $|\nabla I|$  at a fraction of its maximum. This experiment highlights the high-pass nature of naive differencing: while

edges are revealed, amplification of sensor noise and texture is pronounced, motivating the smoothing strategy in the next subsection.


☐  $I_x$  ☐  $I_y$  ☐  $|\nabla I|$  ☒ Edges  $t = 10$




Finite difference — edges ( $t=10$ )

## 1C – Derivative of Gaussian (DoG)

☐ Smooth  $\rightarrow$  Diff ☒ DoG ( $\partial G \otimes I$ ) ☐  $I_x$  ☐  $I_y$  ☒  $|\nabla I|$





1C — DoG,  $|\nabla I|$

I replace raw differences with derivatives of a Gaussian. A normalized Gaussian kernel  $g_\sigma$  is constructed by `proj2/filters.gaussian_kernel(ksize,  $\sigma$ )`, using  $ksize \approx 6\sigma + 1$  to capture sufficient support. Two pipelines are compared: (i) smoothing the image with  $g_\sigma$  and then applying  $\partial x/\partial y$ ; and (ii) forming the DoG filters  $\partial G/\partial x$ ,  $\partial G/\partial y$  (by convolving  $g$  with the difference stencils) and applying a single convolution per axis. Both methods yield near-identical  $I_x$ ,  $I_y$ , and  $|\nabla I|$  when  $\sigma$  and support are matched, but the DoG route is computationally attractive and conceptually cleaner. For the bells-and-whistles visualization, I compute orientations  $\theta = \text{atan2}(I_y, I_x)$  and map them to hue in HSV, with value proportional to  $|\nabla I|$ ; this exposes coherent edge directions across the scene.

## 2A – Unsharp Masking / Sharpening

**Dataset:** Taj ☒ Source ☐ Blurred (low-pass) ☐ High-pass  
☐ Sharpened

2A result

2A — Taj • Source

Unsharp masking is implemented in `proj2/sharpen.py`. A Gaussian blur  $blur = g_\sigma \circledast I$  is subtracted from the input to isolate high frequencies  $high = I - blur$ ; the sharpened output is  $I' = \text{clip}(I + \alpha \cdot high, 0, 1)$ . Parameters  $\sigma$  and  $\alpha$  govern the spatial scale and strength of the enhancement: increasing  $\sigma$  shifts the emphasis to broader features, whereas large  $\alpha$  may introduce halos near strong edges. I also demonstrate reversal by first blurring a sharp

image (setting  $\alpha = -1$  to visualize the low-pass) and then attempting to re-sharpen it; the comparison clarifies that sharpening restores local contrast but cannot recreate frequencies that have been eliminated by the blur.

## 2B – Hybrid Images

**Set:** Example (A+B) ▾ ☐ A (low-freq carrier) ☐ B (high-freq detail) ☐ Low-pass(A)

☐ High-pass(B) ☒ Hybrid ☐ FFT(A) ☐ FFT(B) ☐ FFT(Low) ☐ FFT(High)

☐ FFT(Hybrid)



## 2B — Example • Hybrid

The hybrid construction in `proj2/hybrid.py` forms  $H = LP_{\sigma_L}(A) + HP_{\sigma_H}(B)$ , where  $HP(B) = B - LP_{\sigma_H}(B)$ . The function `hybrid_image(A, B, low_ksize,  $\sigma_L$ , high_ksize,  $\sigma_H$ )` returns the low-pass of  $A$ , the high-pass of  $B$ , their clipped sum, and a log-magnitude FFT for analysis (`log_fft`). Images are pre-aligned and resized via `proj2/io_utils.match_size` so that semantic structures are co-located. Successful hybrids rely on three controls: accurate geometric alignment (to avoid double features), a sufficiently large  $\sigma_L$  that removes mid-frequencies from the carrier, and a moderate  $\sigma_H$  that preserves crisp detail without injecting noise. I tune cutoffs by inspecting the FFTs and by checking whether the interpretation of  $H$  flips with viewing distance as intended.

## 2C – Gaussian & Laplacian Stacks

Image: Apple ▾ ☒ Gaussian Stack ☐ Laplacian Stack




2C — Gaussian stack • Apple

In `stacks.py`, `gaussian_stack(img, L, k,  $\sigma$ )` generates a stack of  $L$  images at full resolution by repeated blurring with the same  $\sigma$ ; `laplacian_stack` stores differences between adjacent Gaussian levels, with the coarsest Gaussian retained as the residual. I visualize stacks using `proj2/viz.save_stack_grid`. Unlike pyramids, stacks avoid down-sampling and simplify subsequent per-pixel operations (e.g., stack-wise masking) at the expense of memory. The Laplacian representation is a band-pass decomposition of the image and is a natural basis for frequency-aware blending.



## 2D – Multiresolution Blending

Mask: Vertical ▾ ☒ Result (Orapple) ☐ Mask Stack ☐ Laplacian Blended Stack



2D — Result (Orapple) • Vertical mask

The blender in `proj2/blend.py` implements the Burt–Adelson recipe using stacks. Given two color images  $A, B$  and a mask  $M \in [0, 1]$ , I build Laplacian stacks  $LA, LB$  and a Gaussian stack of the mask  $GM$ . At each level  $i$ , I compute  $L_{blend}[i] = GM[i] \cdot LA[i] + (1 - GM[i]) \cdot LB[i]$ , then collapse the blended Laplacian stack to reconstruct the output. A hard vertical step mask reproduces the classic “orapple” result; an irregular, feathered mask demonstrates that spatially varying, scale-appropriate smoothing of the seam yields perceptually seamless composites even across textured regions.

JS: 2B viewer